

RAJ
KANE

राज
काणे

[BLOG](#)
[ABOUT](#)
[CONTACT](#)

DEBUGGING STRIPPED BINARIES IN GDB

2022-12-11

Debugging executables is all well and good when they are compiled with GCC's `-g` flag to retain debugging information, but hackers have to deal with **stripped** binaries. By removing symbolic information unessential for correct execution, stripping not only saves disk space and potentially improves performance, but—in the context of security—serves as one level of obfuscation against prying eyes.

We examine the effects of stripping using GCC's `-s` flag on the following toy program that prints out a UID.

```
→ cat getuid.c
#include <stdio.h>
#include <unistd.h>

int main() {
    printf("UID: %d\n", geteuid());
}
→ gcc getuid.c -o getuid
→ ./getuid
UID: 1000
```

After compiling the program without any stripping, we can open the executable with GDB and readily place a breakpoint at a symbol. Shoutout to the extremely handy GDB enhancer, [GEF](#).

```
→ gdb getuid
Reading symbols from getuid...
gef➤ b main
Breakpoint 1 at 0x114d
```

But when we compile and strip the binary, we no longer have this luxury.

```
→ gcc -s getuid.c -o getuid
→ gdb getuid
```

```

Reading symbols from getuid...
Debuginfod has been disabled.
(No debugging symbols found in getuid)
gef> b main
Function "main" not defined.

```

Not to worry! We can still find our way in this seemingly hopeless scenario by finding the program's **entry point** offset, then its **start address**, and finally our destination.

Within GDB, we find the entry point with the command `info files`.

```

gef> info files
Symbols from "/home/raj/getuid".
Local exec file:
  `/home/raj/getuid', file type elf64-x86-64.
Entry point: 0x1050
0x0000000000000318 - 0x0000000000000334 is .interp
0x0000000000000338 - 0x0000000000000378 is .note.gnu.property
0x0000000000000378 - 0x000000000000039c is .note.gnu.build-id
0x000000000000039c - 0x00000000000003bc is .note.ABI-tag
0x00000000000003c0 - 0x00000000000003dc is .gnu.hash
0x00000000000003e0 - 0x00000000000004a0 is .dynsym
0x00000000000004a0 - 0x0000000000000537 is .dynstr
0x0000000000000538 - 0x0000000000000548 is .gnu.version
0x0000000000000548 - 0x0000000000000578 is .gnu.version_r
0x0000000000000578 - 0x0000000000000638 is .rela.dyn
0x0000000000000638 - 0x0000000000000668 is .rela.plt
0x0000000000001000 - 0x000000000000101b is .init
0x0000000000001020 - 0x0000000000001050 is .plt
0x0000000000001050 - 0x000000000000116f is .text
0x0000000000001170 - 0x000000000000117d is .fini
0x0000000000002000 - 0x000000000000200d is .rodata
0x0000000000002010 - 0x0000000000002034 is .eh_frame_hdr
0x0000000000002038 - 0x00000000000020b4 is .eh_frame
0x0000000000003dd0 - 0x0000000000003dd8 is .init_array
0x0000000000003dd8 - 0x0000000000003de0 is .fini_array
0x0000000000003de0 - 0x0000000000003fc0 is .dynamic
0x0000000000003fc0 - 0x0000000000003fe8 is .got
0x0000000000003fe8 - 0x0000000000004010 is .got.plt
0x0000000000004010 - 0x0000000000004020 is .data
0x0000000000004020 - 0x0000000000004028 is .bss

```

Then, we find the start address by doing `set stop-on-solib-events 1`, running the program, and doing `info proc map`.

```

gef> set stop-on-solib-events 1
gef> r
Starting program: /home/raj/getuid
Stopped due to shared library event (no libraries added or removed)
gef> info proc map
process 59703
Mapped address spaces:

```

Start Addr	End Addr	Size	Offset	Perms	objfile
0x555555554000	0x555555555000	0x1000	0x0	r--p	/home/raj/getuid
0x555555555000	0x555555556000	0x1000	0x1000	r-xp	/home/raj/getuid

0x555555556000	0x555555557000	0x1000	0x2000	r--p	/home/raj/getuid
0x555555557000	0x555555559000	0x2000	0x2000	rw-p	/home/raj/getuid
0x7ffff7fc4000	0x7ffff7fc8000	0x4000	0x0	r--p	[vvar]
0x7ffff7fc8000	0x7ffff7fca000	0x2000	0x0	r-xp	[vdso]
0x7ffff7fca000	0x7ffff7fcb000	0x1000	0x0	r--p	/usr/lib/ld-linux-x86-64.so.2
0x7ffff7fcb000	0x7ffff7ffb000	0x26000	0x1000	r-xp	/usr/lib/ld-linux-x86-64.so.2
0x7ffff7ffb000	0x7ffff7ffb000	0xa000	0x27000	r--p	/usr/lib/ld-linux-x86-64.so.2
0x7ffff7ffb000	0x7ffff7fff000	0x4000	0x31000	rw-p	/usr/lib/ld-linux-x86-64.so.2
0x7ffff7fff000	0x7ffff7fff000	0x22000	0x0	rw-p	[stack]
0xffffffff600000	0xffffffff601000	0x1000	0x0	--xp	[vsyscall]

So, we find the two crucial pieces of information that we need: the entry point offset `0x1050` and the start address `0x555555554000`. We set a breakpoint at their sum, and then continue twice to break at the entry point.

```
gef> b *(0x555555554000 + 0x1050)
Breakpoint 1 at 0x555555555050
gef> c
Continuing.
[Thread debugging using libthread_db enabled]
Using host libthread_db library "/usr/lib/libthread_db.so.1".
Stopped due to shared library event:
  Inferior loaded /usr/lib/libc.so.6
gef> c
Continuing.

Breakpoint 1, 0x0000555555555050 in ?? ()
```

We are now at the entry point, but not yet at `main()`. So, we keep pushing and examine the next few instructions.

```
gef> x /15i $rip
=> 0x555555555050:      endbr64
    0x555555555054:      xor     ebp,ebp
    0x555555555056:      mov     r9,rdx
    0x555555555059:      pop     rsi
    0x55555555505a:      mov     rdx,rsi
    0x55555555505d:      and     rsp,0xffffffffffffffff
    0x555555555061:      push    rax
    0x555555555062:      push    rsp
    0x555555555063:      xor     r8d,r8d
    0x555555555066:      xor     ecx,ecx
    0x555555555068:      lea     rdi,[rip+0xda]          # 0x555555555149
    0x55555555506f:      call   QWORD PTR [rip+0x2f4b]  # 0x5555555557fc0
    0x555555555075:      hlt
    0x555555555076:      cs nop WORD PTR [rax+rax*1+0x0]
    0x555555555080:      lea     rdi,[rip+0x2f99]        # 0x5555555558020
```

At instruction `0x55555555506f` we identify the call to `__libc_start_main`, which is located at the address loaded into register `rdi` in the previous instruction. So, we break at `0x555555555149` and continue execution.

```
gef> b *0x555555555149
Breakpoint 2 at 0x555555555149
gef> c
```

Continuing.

Breakpoint 2, 0x000055555555149 in ?? ()

Examining the next few instructions, we see the familiar function prologue and epilogue, as well as markers for the functions in our program.

```
gef> x/15i $rip
=> 0x55555555149:    push    rbp
    0x5555555514a:    mov     rbp, rsp
    0x5555555514d:    call   0x55555555040 <geteuid@plt>
    0x55555555152:    mov     esi, eax
    0x55555555154:    lea     rax, [rip+0xea9]          # 0x555555556004
    0x5555555515b:    mov     rdi, rax
    0x5555555515e:    mov     eax, 0x0
    0x55555555163:    call   0x55555555030 <printf@plt>
    0x55555555168:    mov     eax, 0x0
    0x5555555516d:    pop     rbp
    0x5555555516e:    ret
    0x5555555516f:    add     bl, dh
    0x55555555171:    nop     edx
    0x55555555174:    sub     rsp, 0x8
    0x55555555178:    add     rsp, 0x8
```

Terrific! We finally made it to our `main()` function, and can now begin *actually* reversing (although there isn't much to reverse in our toy UID example).